

- arXiv:2506.05796 [eess.AS]
- [32] Yi Liu, Pascale Fung, Yongsheng Yang, Christopher Cieri, Shudong Huang, and David Graff. 2006. HKUST/MTS: A very large scale Mandarin telephone speech corpus. In *Proceedings of the International Symposium on Chinese Spoken Language Processing (ISCSLP)*. 724–735.
 - [33] Michael McAuliffe, Michaela Socolof, Sarah Mihuc, Michael Wagner, and Morgan Sonderegger. 2017. Montreal Forced Aligner: Trainable Text-Speech Alignment Using Kaldi. In *Proceedings of the Conference of the International Speech Communication Association (Interspeech)*. 498–502.
 - [34] Ivan Medennikov, Maxim Korenevsky, Tatiana Prisyach, Yuri Khokhlov, et al. 2020. Target-Speaker Voice Activity Detection: A Novel Approach for Multi-Speaker Diarization in a Dinner Party Scenario. In *Proceedings of the Conference of the International Speech Communication Association (Interspeech)*. 274–278.
 - [35] Lingwei Meng, Shujie Hu, Jiawen Kang, Zhaoqing Li, Yuejiao Wang, Wenxuan Wu, Xixin Wu, Xunying Liu, and Helen Meng. 2025. Large Language Model Can Transcribe Speech in Multi-Talker Scenarios with Versatile Instructions. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 1–5.
 - [36] Lingwei Meng, Shujie Hu, Jiawen Kang, Zhaoqing Li, Yuejiao Wang, Wenxuan Wu, Xixin Wu, Xunying Liu, and Helen Meng. 2025. Large language model can transcribe speech in multi-talker scenarios with versatile instructions. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 1–5.
 - [37] Bingshen Mu, Pengcheng Guo, Zhaokai Sun, Shuai Wang, Hexin Liu, Mingchen Shao, Lei Xie, Eng Siong Chng, Longshuai Xiao, Qiangze Feng, et al. 2025. Summary on The Multilingual Conversational Speech Language Model Challenge: Datasets, Tasks, Baselines, and Methods. *arXiv preprint arXiv:2509.13785* (2025).
 - [38] Douglas O'Shaughnessy. 2025. Speaker Diarization: A Review of Objectives and Methods. *Applied Sciences* 15, 4 (2025), 2002.
 - [39] Taejin Park, Ivan Medennikov, Kunal Dhawan, Weiqing Wang, He Huang, Nithin Rao Koluguri, Krishna C Puvvada, Jagadeesh Balam, and Boris Ginsburg. 2025. Sortformer: A Novel Approach for Permutation-Resolved Speaker Supervision in Speech-to-Text Systems. In *Proceedings of the Forty-second International Conference on Machine Learning*.
 - [40] Tae Jin Park, Kunal Dhawan, Nithin Koluguri, and Jagadeesh Balam. 2024. Enhancing speaker diarization with large language models: A contextual beam search approach. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 10861–10865.
 - [41] Tae Jin Park, Naoyuki Kanda, Dimitrios Dimitriadis, Kyu J. Han, Shinji Watanabe, et al. 2022. A review of speaker diarization: Recent advances with deep learning. *Computer Speech & Language* 72 (2022).
 - [42] Zhiliang Peng, Jianwei Yu, Yaoyao Chang, Zilong Wang, Li Dong, Yingbo Hao, Yujie Tu, Chenyu Yang, Wenhui Wang, Songchen Xu, Yutao Sun, Hangbo Bao, Weijiang Xu, Yi Zhu, Zehua Wang, Ting Song, Yan Xia, Zewen Chi, Shaohan Huang, Liang Wang, Chuang Ding, Shuai Wang, Xie Chen, and Furu Wei. 2026. VIBEVOICE-ASR Technical Report. *arXiv:2601.18184* [cs.SD]
 - [43] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. 2023. Robust speech recognition via large-scale weak supervision. In *Proceedings of the International conference on machine learning*. PMLR, 28492–28518.
 - [44] Desh Raj, Pavel Denisov, Zhuo Chen, Hakan Erdogan, Zili Huang, Maokui He, Shinji Watanabe, Jun Du, Takuya Yoshioka, Yi Luo, et al. 2021. Integration of speech separation, diarization, and recognition for multi-speaker meetings: System description, comparison, and analysis. In *Proceedings of the spoken language technology workshop (SLT)*. IEEE, 897–904.
 - [45] Mohan Shi, Xiong Xiao, Ruchao Fan, Shaoshi Ling, and Jinyu Li. 2025. Train Short, Infer Long: Speech-LLM Enables Zero-Shot Streamable Joint ASR and Diarization on Long Audio. *arXiv preprint arXiv:2511.16046* (2025).
 - [46] Jinhan Wang, Weiqing Wang, Kunal Dhawan, Taejin Park, Myungjong Kim, Ivan Medennikov, He Huang, Nithin Koluguri, Jagadeesh Balam, and Boris Ginsburg. 2025. META-CAT: Speaker-Informed Speech Embeddings via Meta Information Concatenation for Multi-talker ASR. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 1–5.
 - [47] Quan Wang, Carlton Downey, Li Wan, Mansfield, et al. 2018. Speaker Diarization with LSTM. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 5239–5243.
 - [48] Quan Wang, Yiling Huang, Guanlong Zhao, Evan Clark, Wei Xia, and Hank Liao. 2024. DiarizationLM: Speaker diarization post-processing with large language models. *arXiv preprint arXiv:2401.03506* (2024).
 - [49] Weiqing Wang, Danwei Cai, Qingjian Lin, et al. 2021. The dku-dukeeece-lenovo system for the diarization task of the 2021 voxceleb speaker recognition challenge. *arXiv preprint arXiv:2109.02002* (2021).
 - [50] Weiqing Wang, Qingjian Lin, Danwei Cai, and Ming Li. 2022. Similarity Measurement of Segment-Level Speaker Embeddings in Speaker Diarization. *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 30 (2022), 2645–2658.
 - [51] Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. 2022. Emergent Abilities of Large Language Models. *Transactions on Machine Learning Research* (2022).
 - [52] Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, Jialin Wang, Yang Fan, Kai Dang, et al. 2025. Qwen2.5-omni technical report. *arXiv preprint arXiv:2503.20215* (2025).
 - [53] Zehui Yang, Yifan Chen, Lei Luo, Runyan Yang, Lingxuan Ye, Gaofeng Cheng, Ji Xu, Yaohui Jin, Qingqing Zhang, Pengyuan Zhang, Lei Xie, and Yonghong Yan. 2022. Open Source MagicData-RAMC: A Rich Annotated Mandarin Conversational (RAMC) Speech Dataset. In *Proceedings of the Conference of the International Speech Communication Association (Interspeech)*.
 - [54] Han Yin, Yafeng Chen, Chong Deng, Luyao Cheng, Hui Wang, Chao-Hong Tan, Qian Chen, Wen Wang, and Xiangang Li. 2025. SpeakerLM: End-to-end versatile speaker diarization and recognition with multimodal large language models. *arXiv preprint arXiv:2508.06372* (2025).
 - [55] Donghua Yu, Zhengyuan Lin, Chen Yang, Yiyang Zhang, Zhaoye Fei, Hanfu Chen, Jingqi Chen, Ke Chen, Qinyuan Cheng, Liwei Fan, et al. 2026. MOSS Transcribe Diarize: Accurate Transcription with Speaker Diarization. *arXiv preprint arXiv:2601.01554* (2026).
 - [56] F. Yu, S. Zhang, Y. Fu, L. Xie, S. Zheng, Z. Du, et al. 2022. M2met: The icassp 2022 multi-channel multi-party meeting transcription challenge. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 6167–6171.
 - [57] Fan Yu, Shiliang Zhang, Pengcheng Guo, Yihui Fu, Zhihao Du, Siqi Zheng, Weilong Huang, Lei Xie, Zheng-Hua Tan, DeLiang Wang, et al. 2022. Summary on the ICASSP 2022 multi-channel multi-party meeting transcription grand challenge. In *Proceedings of the International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 9156–9160.
 - [58] W. Zhang, X. Chang, Y. Qian, and S. Watanabe. 2020. Improving end-to-end single-channel multi-talker speech recognition. *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 28 (2020), 1385–1394.
 - [59] Xianrui Zheng, Chao Zhang, and Phil Woodland. 2025. DNCASR: End-to-End Training for Speaker-Attributed ASR. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 18369–18383.